



Entity Relationship Model

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Design Phases

- Initial phase -- characterize fully the data needs of the prospective database users.
- Second phase -- choosing a data model
 - Applying the concepts of the chosen data model
 - Translating these requirements into a conceptual schema of the database.
 - A fully developed conceptual schema indicates the functional requirements of the enterprise.
 - Describe the kinds of operations (or transactions) that will be performed on the data.



Design Phases (Cont.)

- Final Phase -- Moving from an abstract data model to the implementation of the database
 - Logical Design – Deciding on the database schema.
 - Database design requires that we find a “good” collection of relation schemas.
 - Business decision – What attributes should we record in the database?
 - Computer Science decision – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
 - Physical Design – Deciding on the physical layout of the database



Design Approaches

- Entity Relationship Model
 - Models an enterprise as a collection of *entities* and *relationships*
 - Entity: a “thing” or “object” in the enterprise that is distinguishable from other objects
 - Described by a set of *attributes*
 - Relationship: an association among several entities
 - Represented diagrammatically by an *entity-relationship diagram*



Entity Sets

- An **entity** is an object that exists and is distinguishable from other objects.
 - Example: specific person, company, event, plant
- An **entity set** is a set of entities of the same type that share the same properties.
 - Example: set of all persons, companies, trees, holidays
- An entity is represented by a set of attributes; i.e., descriptive properties possessed by all members of an entity set.
 - Example:
 $\text{instructor} = (\text{ID}, \text{name}, \text{salary})$
 $\text{course} = (\text{course_id}, \text{title}, \text{credits})$
- A subset of the attributes form a **primary key** of the entity set; i.e., uniquely identifying each member of the set.



Representing Entity sets in ER Diagram

- Entity sets can be represented graphically as follows:
 - Rectangles represent entity sets.
 - Attributes listed inside entity rectangle
 - Underline indicates primary key attributes

<i>instructor</i>
<u>ID</u>
<i>name</i>
<i>salary</i>

<i>student</i>
<u>ID</u>
<i>name</i>
<i>tot_cred</i>



Weak Entity Sets

- Consider a *section* entity, which is uniquely identified by a *course_id*, *semester*, *year*, and *sec_id*.
- Clearly, section entities are related to course entities. Suppose we create a relationship set *sec_course* between entity sets *section* and *course*.
- Note that the information in *sec_course* is redundant, since *section* already has an attribute *course_id*, which identifies the course with which the section is related.
- One option to deal with this redundancy is to get rid of the relationship *sec_course*; however, by doing so the relationship between *section* and *course* becomes implicit in an attribute, which is not desirable.



Weak Entity Sets (Cont.)

- An alternative way to deal with this redundancy is to not store the attribute *course_id* in the *section* entity and to only store the remaining attributes *section_id*, *year*, and *semester*.
 - However, the entity set *section* then does not have enough attributes to identify a particular *section* entity uniquely
- To deal with this problem, we treat the relationship *sec_course* as a special relationship that provides extra information, in this case, the *course_id*, required to identify *section* entities uniquely.
- A **weak entity set** is one whose existence is dependent on another entity, called its **identifying entity**
- Instead of associating a primary key with a weak entity, we use the identifying entity, along with extra attributes called **discriminator** to uniquely identify a weak entity.



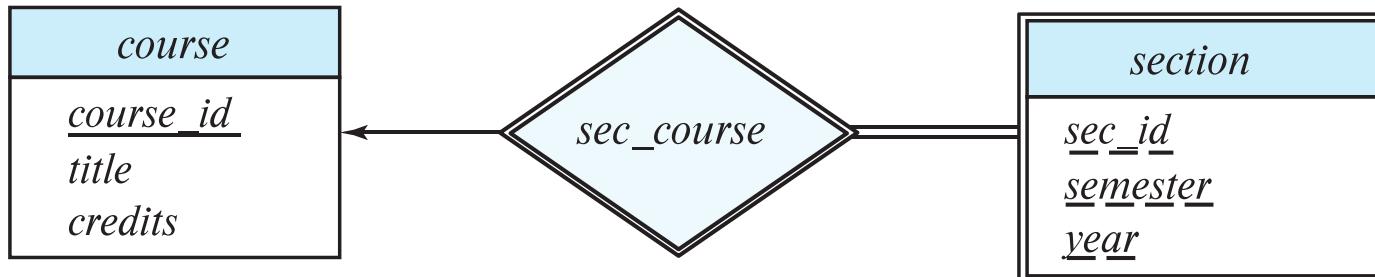
Weak Entity Sets (Cont.)

- An entity set that is not a weak entity set is termed a **strong entity set**.
- Every weak entity must be associated with an identifying entity; that is, the weak entity set is said to be **existence dependent** on the identifying entity set.
- The identifying entity set is said to **own** the weak entity set that it identifies.
- The relationship associating the weak entity set with the identifying entity set is called the **identifying relationship**.



Expressing Weak Entity Sets

- In E-R diagrams, a weak entity set is depicted via a double rectangle.
- We underline the discriminator of a weak entity set with a dashed line.
- The relationship set connecting the weak entity set to the identifying strong entity set is depicted by a double diamond.
- Primary key for *section* – (*course_id*, *sec_id*, *semester*, *year*)





Relationship Sets

- A **relationship** is an association among several entities

Example:

44553 (<u>Priti</u>)	<u>advisor</u>	22222 (<u>Mukesh</u>)
<i>student</i> entity	relationship set	<i>instructor</i> entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where (e_1, e_2, \dots, e_n) is a relationship

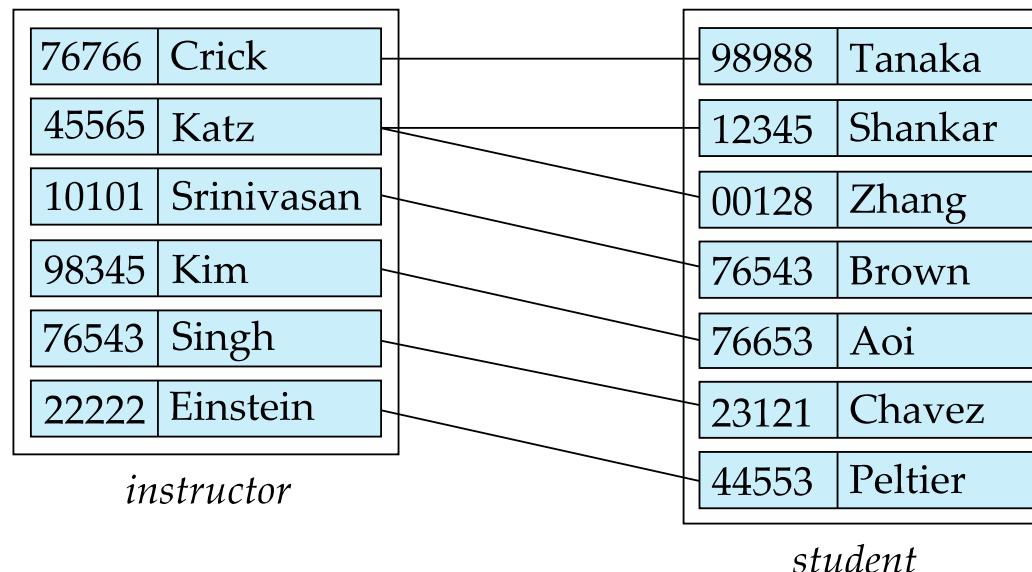
- Example:

$$(44553, 22222) \in \text{advisor}$$



Relationship Sets (Cont.)

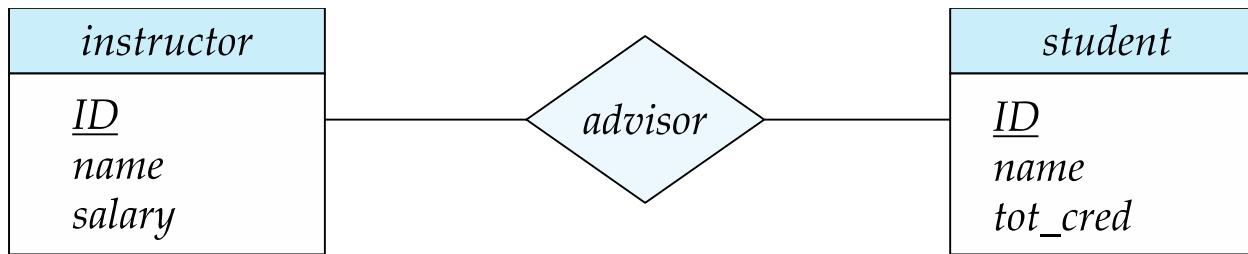
- Example: we define the relationship set *advisor* to denote the associations between students and the instructors who act as their advisors.
- Pictorially, we draw a line between related entities.





Representing Relationship Sets via ER Diagrams

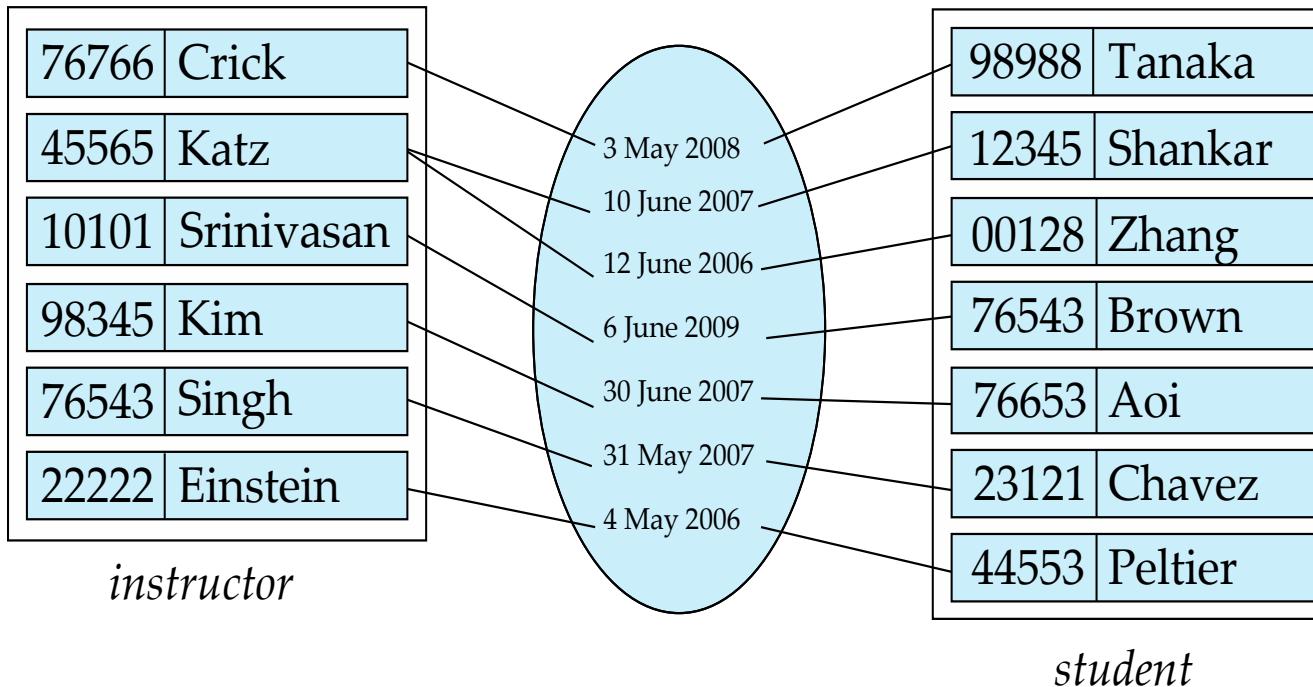
- Diamond shape represent relationship sets.





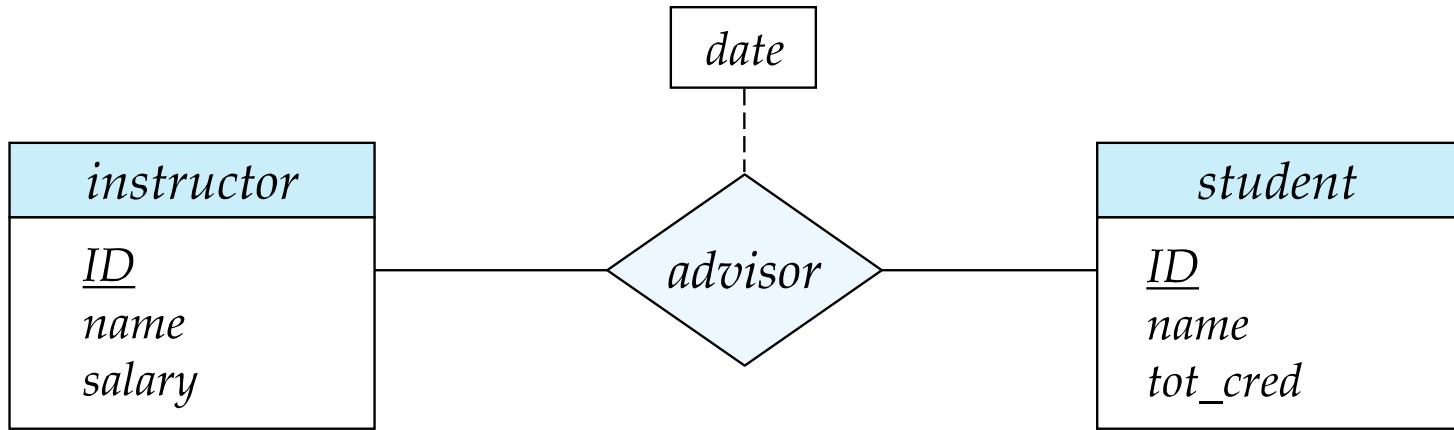
Relationship Sets (Cont.)

- An attribute can also be associated with a relationship set.
- For instance, the *advisor* relationship set between entity sets *instructor* and *student* may have the attribute *date* which tracks when the student started being associated with the advisor





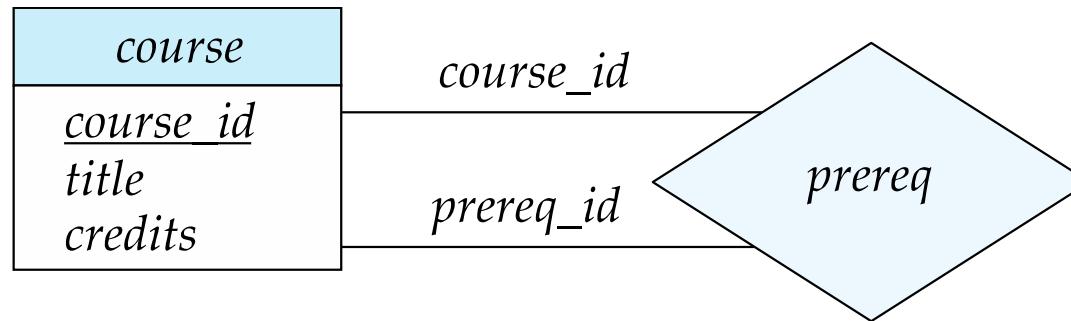
Relationship Sets with Attributes





Roles

- Entity sets of a relationship need not be distinct
 - Each occurrence of an entity set plays a “role” in the relationship
- The labels “*course_id*” and “*prereq_id*” are called **roles**.





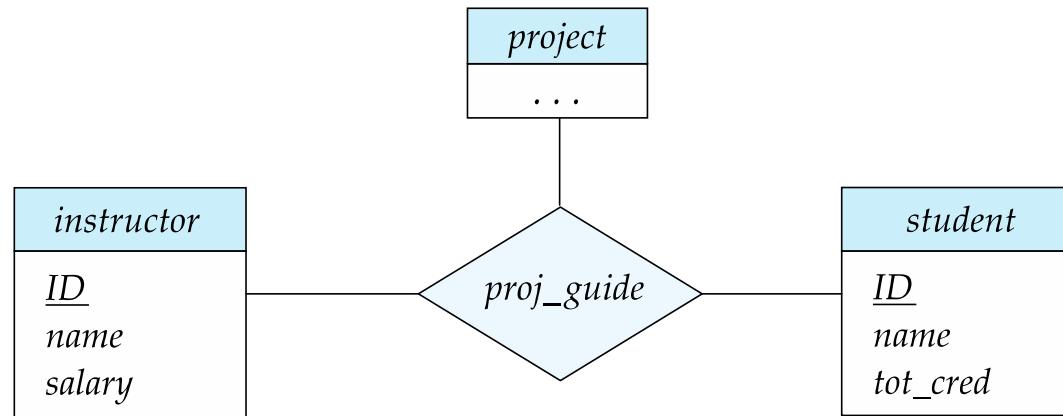
Degree of a Relationship Set

- Binary relationship
 - involve two entity sets (or degree two).
 - most relationship sets in a database system are binary.
- Relationships between more than two entity sets are rare. Most relationships are binary.
 - Example: *students* work on research *projects* under the guidance of an *instructor*.
 - relationship *proj_guide* is a ternary relationship between *instructor*, *student*, and *project*



Non-binary Relationship Sets

- There are occasions when it is more convenient to represent relationships as non-binary.
- E-R Diagram with a Ternary Relationship





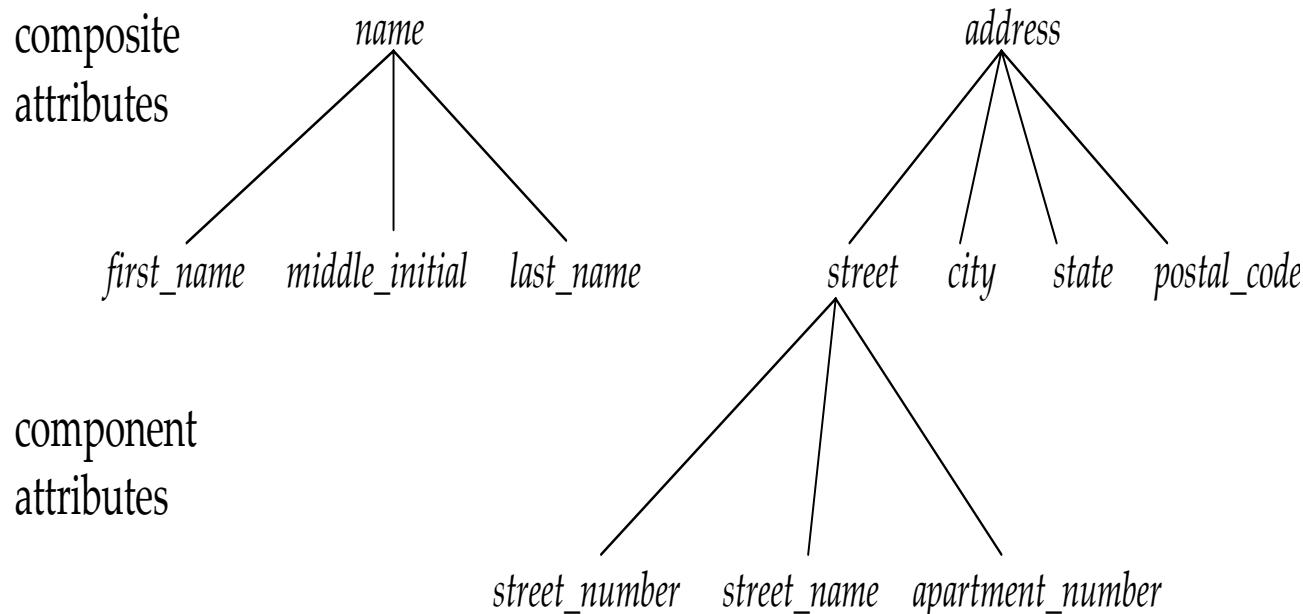
Complex Attributes

- Attribute types:
 - **Simple** and **composite** attributes.
 - **Single-valued** and **multivalued** attributes
 - Example: multivalued attribute: *phone_numbers*
 - **Derived** attributes
 - Can be computed from other attributes
 - Example: age, given date_of_birth
- **Domain** – the set of permitted values for each attribute



Composite Attributes

- Composite attributes allow us to divide attributes into subparts (other attributes).





Representing Complex Attributes in ER Diagram

1. Simple Attribute : Stored as atomic values in a single column

2. Composite Attribute: Stored as separate columns for each component or as a structured data type

3. Multi-Valued: Stored in a separate table with foreign key reference to the main entity

4. Derived Attribute: Not stored directly; calculated at query time

5. Complex Attribute(Combination of composite and multi-valued attributes): Requires nested tables

<i>instructor</i>
<i>ID</i>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age ()</i>

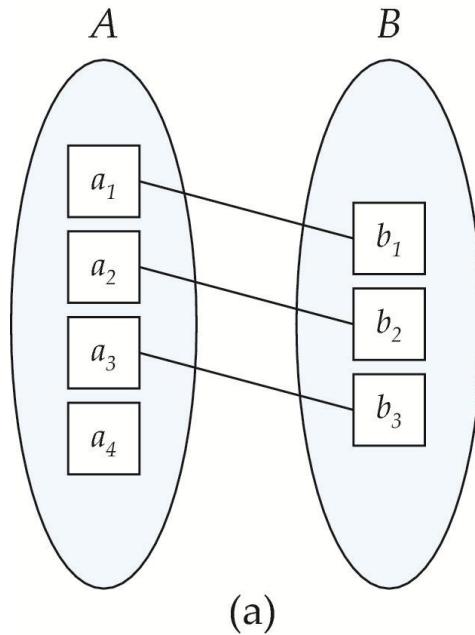


Mapping Cardinality Constraints

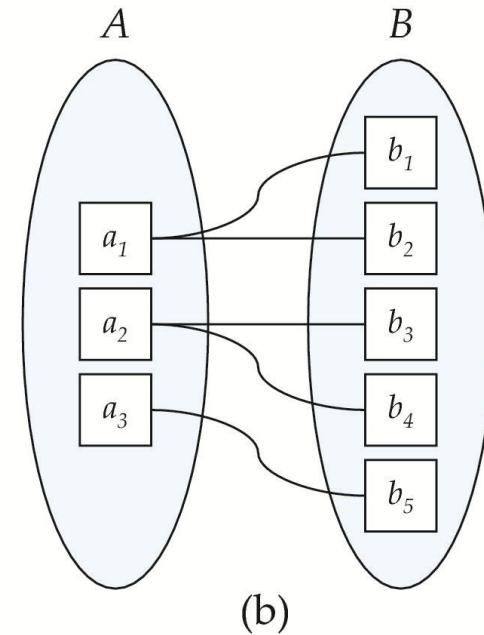
- Express the number of entities to which another entity can be associated via a relationship set.
- Most useful in describing binary relationship sets.
- For a binary relationship set the mapping cardinality must be one of the following types:
 - One to one
 - One to many
 - Many to one
 - Many to many



Mapping Cardinalities



One to one

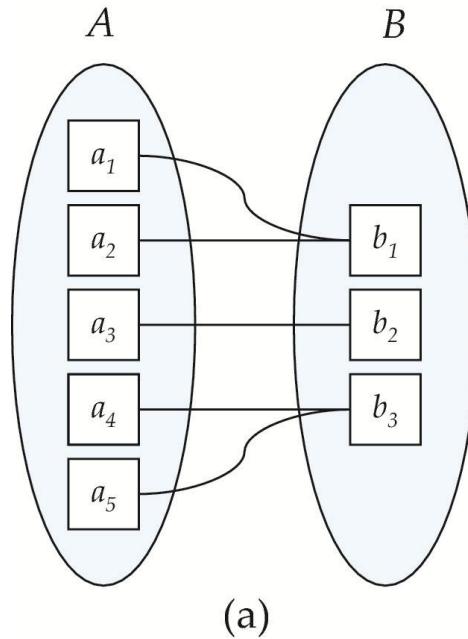


One to many

Note: Some elements in A and B may not be mapped to any elements in the other set

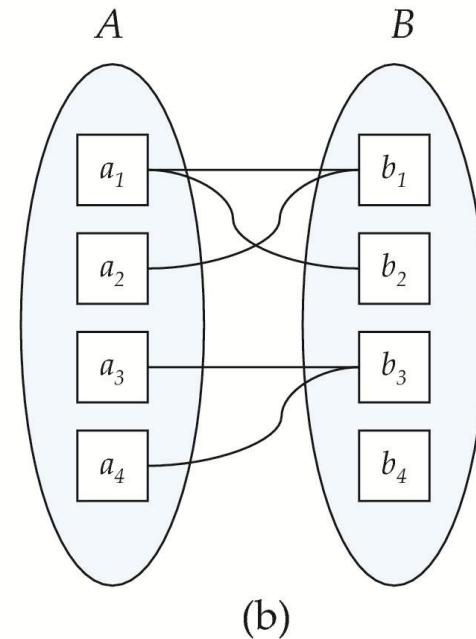


Mapping Cardinalities



(a)

Many to one



(b)

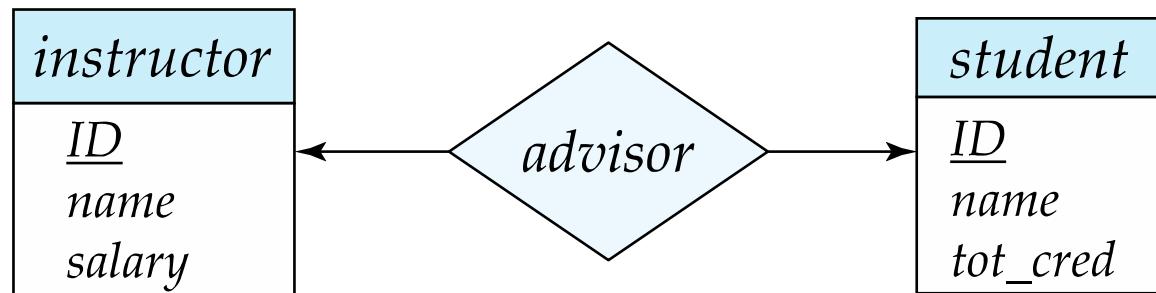
Many to many

Note: Some elements in A and B may not be mapped to any elements in the other set



Representing Cardinality Constraints in ER Diagram

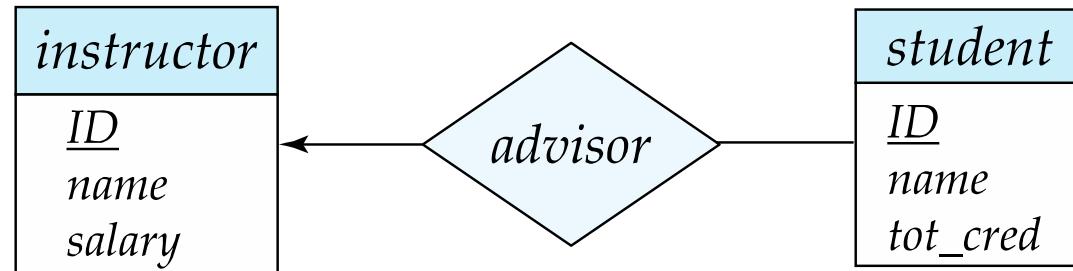
- We express cardinality constraints by drawing either a directed line (\rightarrow), signifying “one,” or an undirected line (—), signifying “many,” between the relationship set and the entity set.
- One-to-one relationship between an *instructor* and a *student* :
 - A *student* is associated with at most one *instructor* via the relationship *advisor*
 - A *student* is associated with at most one *department* via *stud_dept*





One-to-Many Relationship

- one-to-many relationship between an *instructor* and a *student*
 - an instructor is associated with several (including 0) students via *advisor*
 - a student is associated with at most one instructor via advisor,





Many-to-One Relationships

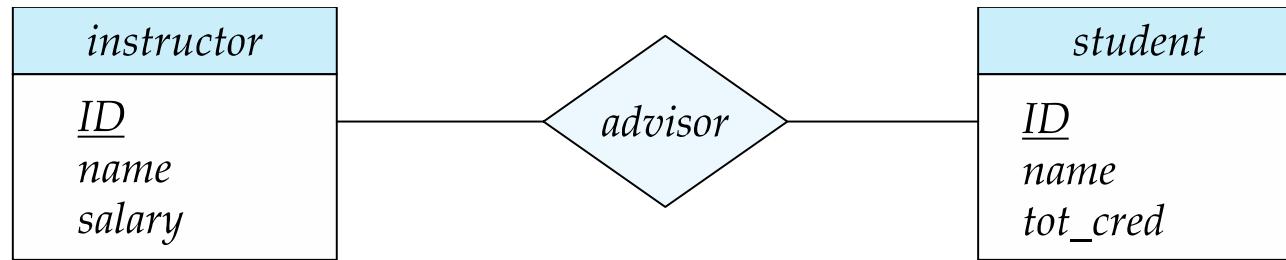
- In a many-to-one relationship between an *instructor* and a *student*,
 - an *instructor* is associated with at most one *student* via *advisor*,
 - and a *student* is associated with several (including 0) *instructors* via *advisor*





Many-to-Many Relationship

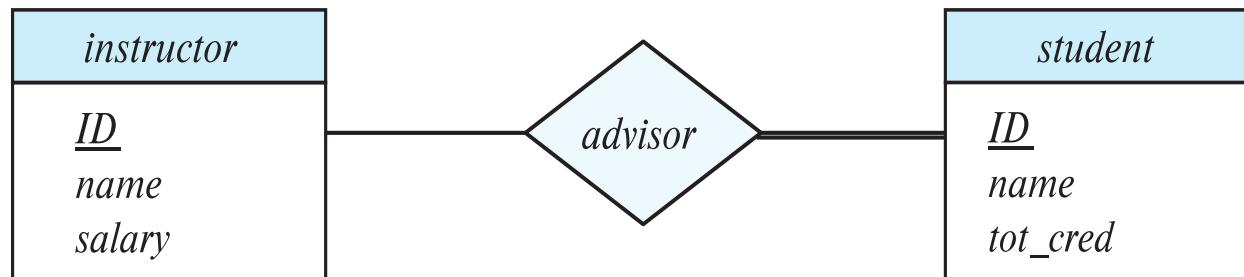
- An instructor is associated with several (possibly 0) students via *advisor*
- A student is associated with several (possibly 0) instructors via *advisor*





Total and Partial Participation

- **Total participation** (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set



participation of *student* in *advisor* relation is total

- every *student* must have an associated *instructor*
- **Partial participation:** some entities may not participate in any relationship in the relationship set
 - Example: participation of *instructor* in *advisor* is partial



Keys

- Used to specify how entities within a given entity set are distinguished
- In other words, no two entities in an entity set are allowed to have exactly the same value for all attributes
- Keys are applicable to relation schemas also **to identify relationships uniquely**, and thus distinguish relationships from each other

Types:

- Super key
- Candidate key
- Primary key
- Alternate Key
- Foreign key



Keys

- Let $K \subseteq R$
- K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$
 - Example: $\{ID\}$ and $\{ID, name\}$ are both superkeys of *instructor*.
- Superkey K is a **candidate key** if K is minimal
Example: $\{ID\}$ is a candidate key for *Instructor*
- One of the candidate keys is selected to be the **primary key**. Others are said to be **Alternate key**.
 - $\{ID\} \rightarrow$ Primary Key
- **Foreign key** constraint: Value in one relation must appear in another
 - **Referencing** relation
 - Example – $dept_name$ in *instructor* is a foreign key from *instructor* referencing *department*



Keys

- Used to uniquely identify any record or row of data from the table.
- Used to establish and identify relationships between tables.
- **For example:**
 - Student table, ID is used as a key because it is unique for each student.
 - PERSON table, passport_number, license_number, SSN are keys since they are unique for each person.

STUDENT
ID
Name
Address
Course

PERSON
Name
DOB
Passport_Number
License_Number
SSN



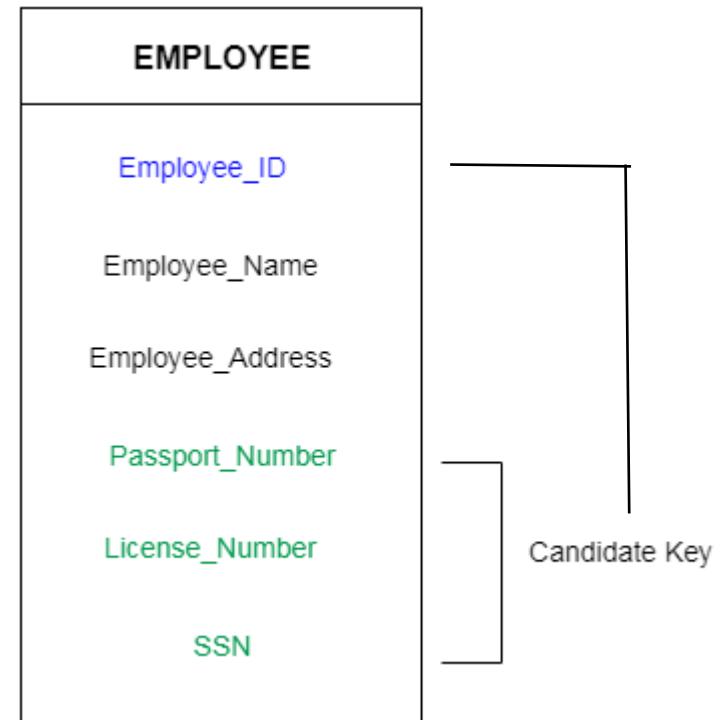
Super Key

- A set of attributes which can uniquely identify a tuple.
- Super key is a superset of a candidate key.
- EMPLOYEE table,
 - (EMPLOYEE_ID, EMPLOYEE_NAME) the name of two employees can be the same, but their EMPLOYEE_ID can't be the same. Hence, this combination can also be a key.
 - The super key would be EMPLOYEE-ID, (EMPLOYEE_ID, EMPLOYEE-NAME), etc.



Candidate Key

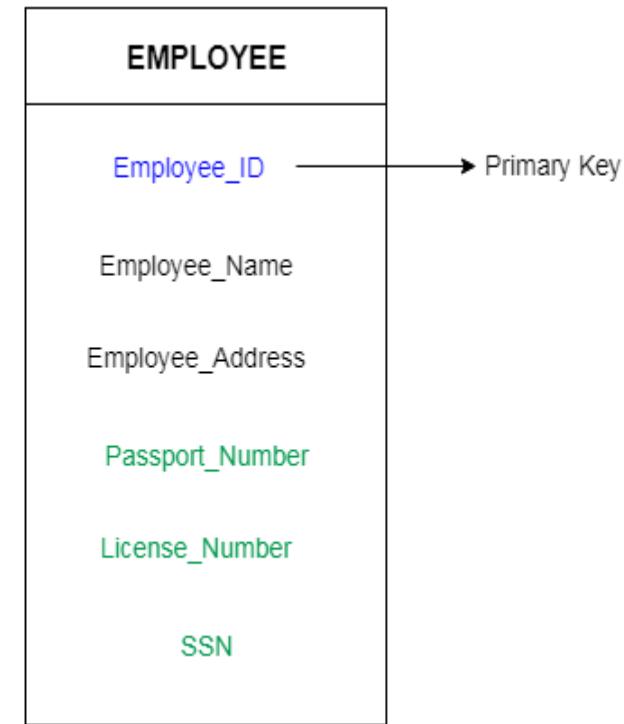
- A candidate key is an attribute or set of attributes which can uniquely and minimally identify a tuple.
- Thus, from the number of obtained candidate keys, we can identify the appropriate primary key.
- E.g. EMPLOYEE table:
 - {employee_id}, {Pass_no}, {L_num},{SSN}





Primary Key

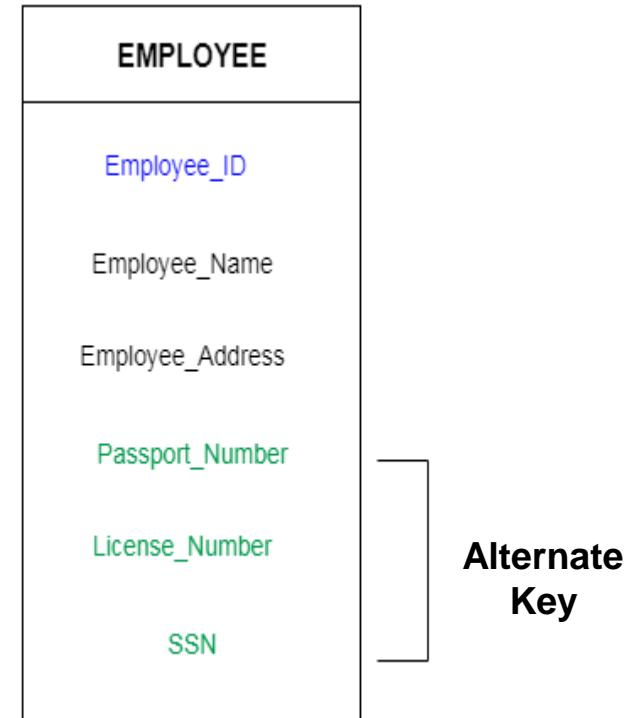
- The first key which is used to identify one and only one instance of an entity uniquely.
- An entity can contain multiple keys
- The key which is most suitable from those lists become a primary key.
- E.g. EMPLOYEE table, ID can be primary key since it is unique for each employee.
License_Number and Passport_Number as primary key since they are also unique.
- For each entity, selection of the primary key is based on requirement and developers.





Alternate Key

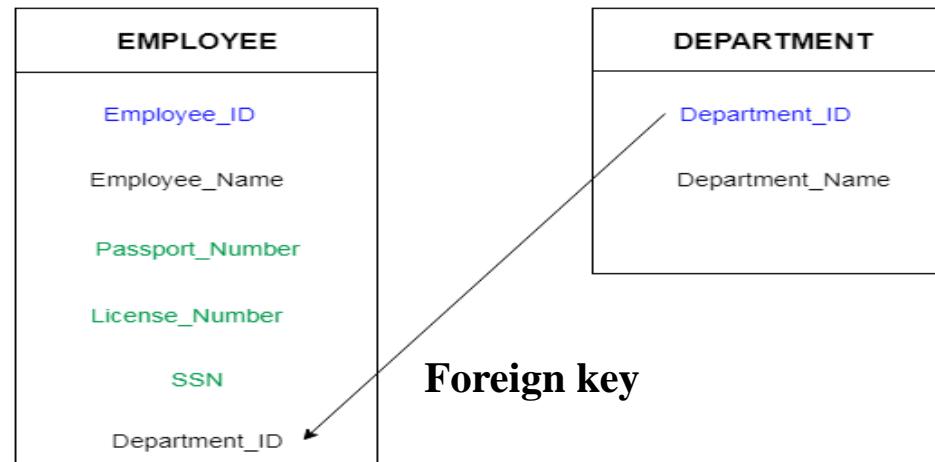
- The alternative key is the key in the database table which is not selected as a primary key.
- It's an additional key similar to the primary key with which we can access data in the table.
- It can uniquely identify attributes of each element in the row of the table.
- But, because there can be multiple candidate keys all cannot be selected. So, the remaining keys which are not selected as primary keys are called alternative keys.
- E.g. EMPLOYEE table:
 - {employee_id}, {Pass_no}, {L_num},{SSN}
 - employee_id can be primary key since it is unique for each employee.
 - {Pass_no}, {L_num},{SSN} they are also unique but not selected as primary key
 - {Pass_no}, {L_num},{SSN} are Alternate key





Foreign Key

- The column of the table which is used to point to the primary key of another table.
 - E.g. In a company, every employee works in a specific department, and employee and department are two different entities.
 - The information of the Employee and Department are in the employee table.
 - Link these two tables through the primary key of one table.
- Steps:
 - Add the primary key of the DEPARTMENT table, Department_Id as a new attribute in the EMPLOYEE table.
 - Now in the EMPLOYEE table, Department_Id is the foreign key, and both the tables are related.





Student Table

Student_Number	Student_Name	Student_Phone	Subject_Number
1	Andrew	6615927284	10
2	Sara	6583654865	20
3	Harry	4647567463	10

- **The Super Keys –**

{ Student_Number } { Student_Phone } { Student_Number,Student_Name }
{ Student_Number,Student_Phone } { Student_Number,Subject_Number }
{ Student_Phone,Student_Name } { Student_Phone,Subject_Number }
{ Student_Number,Student_Name,Student_Phone }
{ Student_Number,Student_Phone,Subject_Number }
{ Student_Number,Student_Name,Subject_Number }
{ Student_Phone,Student_Name,Subject_Number }

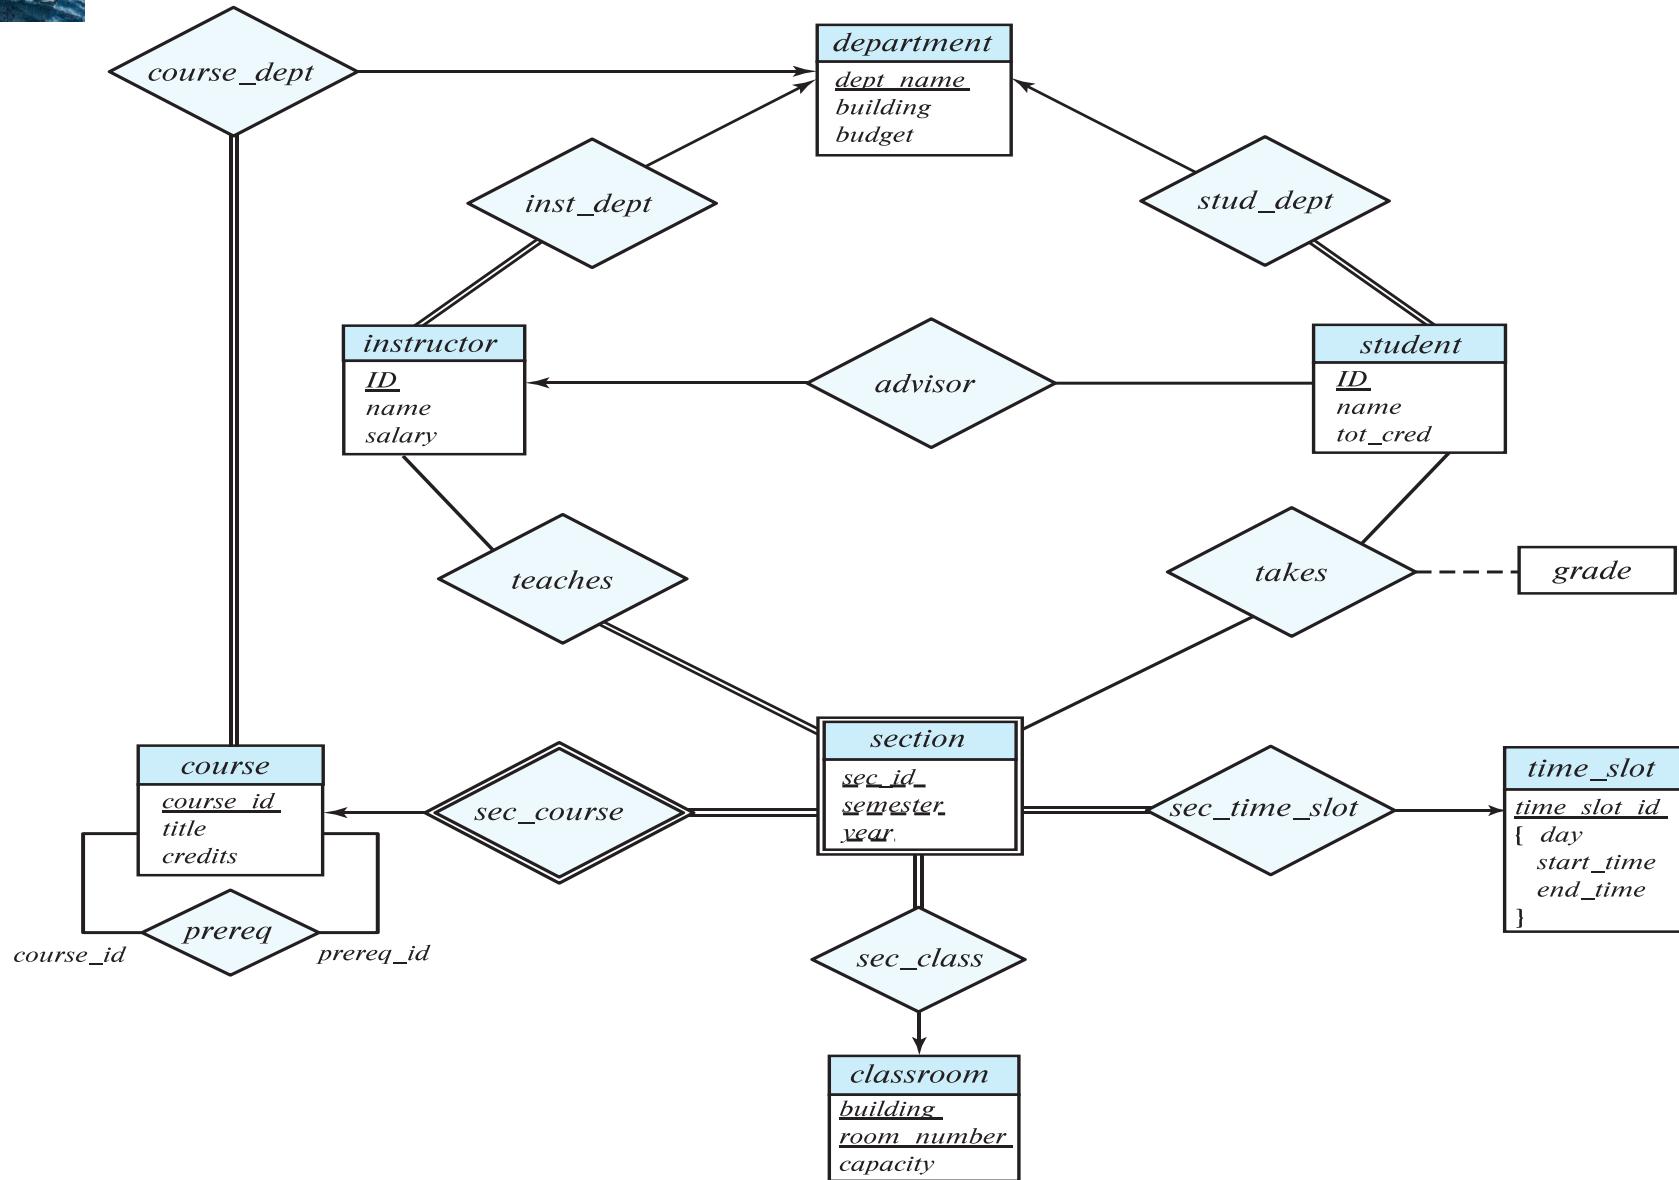
- **The Candidate Key**

{ Student_Number }
{ Student_Phone }
{ Student_name, Subject_Number }

- **The Primary Key : { Student_Number }**



E-R Diagram for a University Enterprise





Extended E-R Features: Specialization

- Top-down design process
- Sub grouping within an entity set that are distinctive from other entities in the set
- These subgroupings become lower-level entity sets that have attributes or participate in relationships that do not apply to the higher-level entity set
- **Attribute inheritance** – a lower-level entity set inherits all the attributes and relationship participation of the higher-level entity set to which it is linked

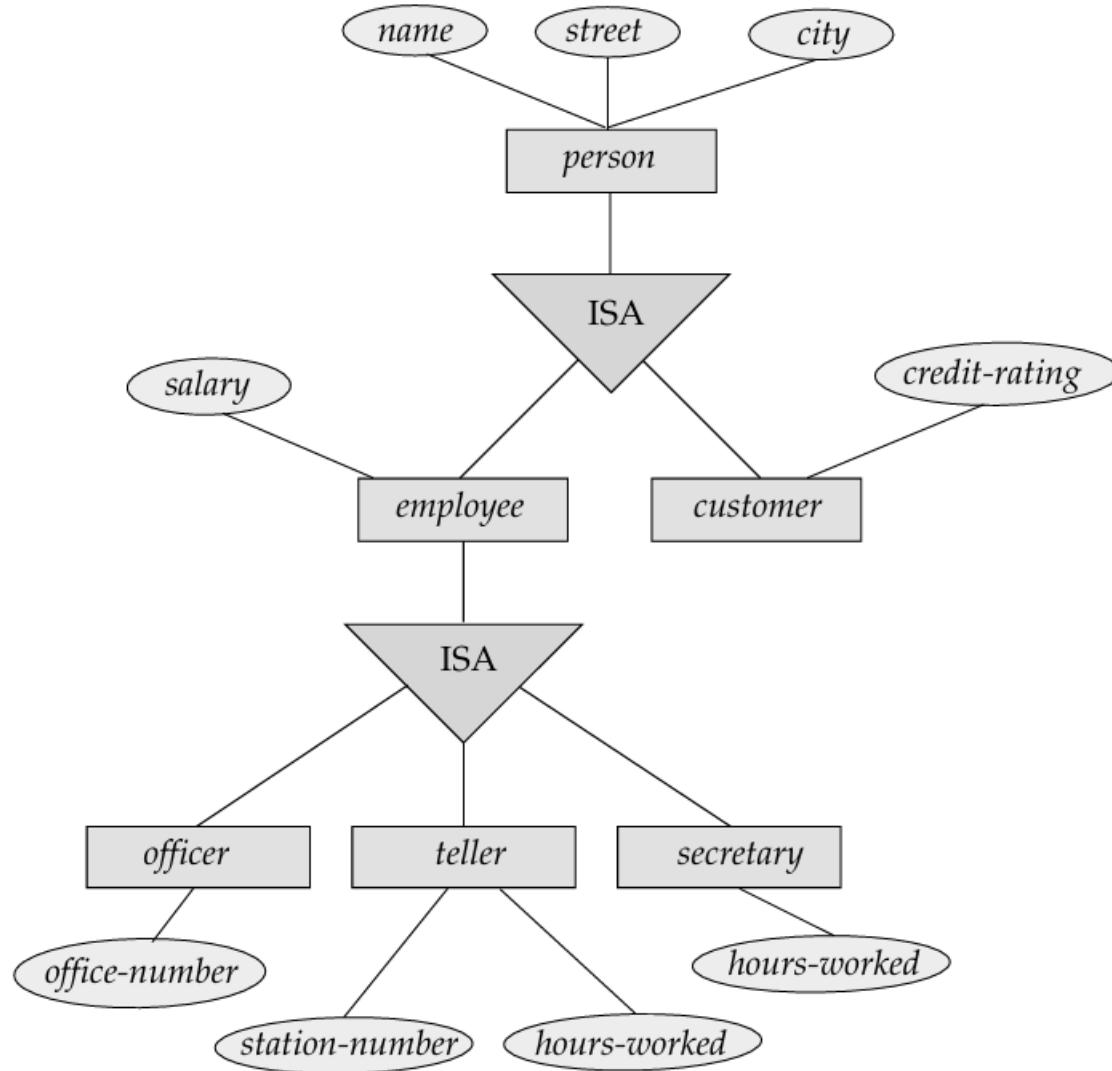


Extended ER Features: Generalization

- A **bottom-up design process** – combine a number of entity sets that share the same features into a higher-level entity set
- Specialization and generalization are simple inversions of each other; they are represented in an E-R diagram in the same way
- The terms specialization and generalization are used interchangeably



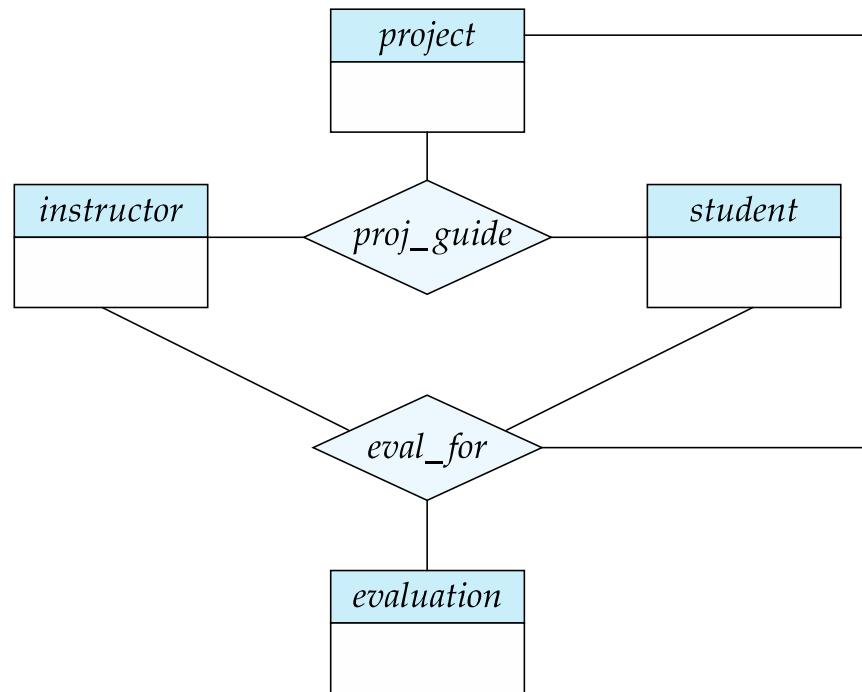
Specialization Example





Aggregation

- Consider the ternary relationship *proj_guide*, which we saw earlier
- Suppose we want to record evaluations of a student by a guide on a project





Aggregation (Cont.)

- Relationship sets *eval_for* and *proj_guide* represent overlapping information
 - Every *eval_for* relationship corresponds to a *proj_guide* relationship
 - However, some *proj_guide* relationships may not correspond to any *eval_for* relationships
 - So we can't discard the *proj_guide* relationship
- Eliminate this redundancy via *aggregation*
 - Treat relationship as an abstract entity
 - Allows relationships between relationships
 - Abstraction of relationship into new entity



Aggregation (Cont.)

- Eliminate this redundancy via *aggregation* without introducing redundancy, the following diagram represents:
 - A student is guided by a particular instructor on a particular project
 - A student, instructor, project combination may have an associated evaluation

